

ORES Preparation IV

Tom Gülenman

17. Dezember 2018

Disclaimer: No guarantee for the correctness of information / explanations / sources is given.

Goals

1. Adjust crucial metrics list
 - to match the damaging model metrics ✓
 - add 2 last metrics
2. Check out mail attachments ✓
3. Check out new Confluence pages and goals ✓
4. Research
 - Check out FAT Conference Docs
 - In what other cases than confusion matrices are those parameters explained?
 - Are there already visualizations of some of these parameters in any contexts?
 - Are there any applications, where I can filter for these parameters → visualizations or just about anything?

1 Crucial metrics: damaging-model

Metrics simple list:

!f1	✓
!precision	✓
!recall	✓
accuracy	✓
counts	
f1	✓
filter_rate	✓
fpr	✓
match_rate	✓
pr_auc	✓
precision	✓
rates	
recall	✓
roc_auc	✓

The metrics are the same for the damaging and itemquality models, but a few changes will be made to the explanatory parts nonetheless (compared to the version in oresDoc3). Also the structure of explanations will be changed to the following:

For each metric (if possible) there will be:


1. An intuitive explanation
2. The formula based on the **confusion matrix**
3. Its meaning based on the “**confusion circle**”
4. Its meaning based on the **loan threshold** representation by Google (Link)

Explanations: References

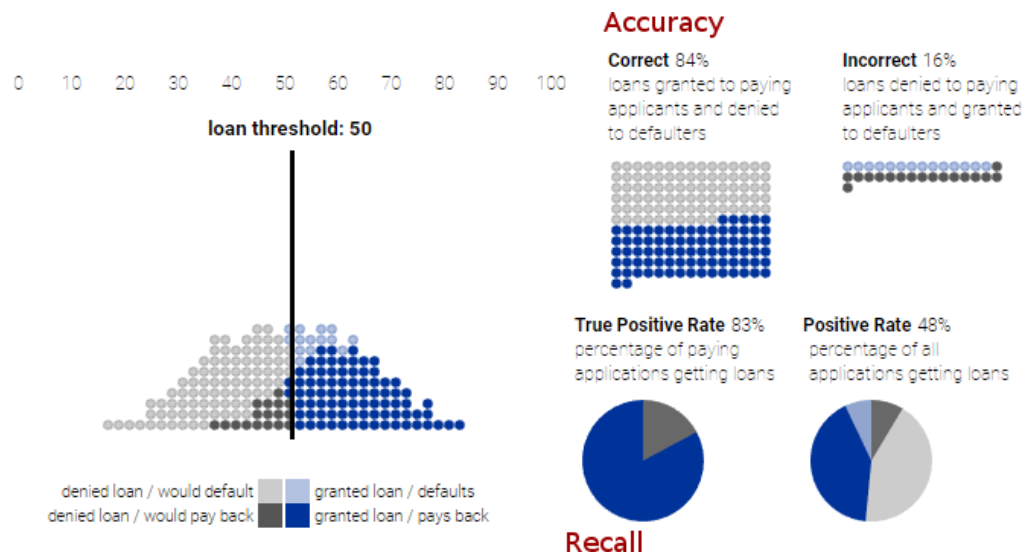
- Confusion Matrix

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

- “Confusion Circle”




- Loan Threshold




1.1 recall

1. Recall (\equiv True Positive Rate) is defined as the ability of a model to find all relevant cases within the dataset.



1.2 precision

1. Ability of the model to find only relevant cases within the dataset



1.3 f1

1. F1-Score, the harmonic mean of recall and precision, a metric from 0 (worst) to 1 (best), used to evaluate the accuracy of a model by taking recall and precision into account


2.	3.	4.
-	-	-

$$= 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

Compared to the simple average (of recall and precision), the harmonic mean punishes extreme values (e.g. precision 1.0 and recall 0.0 \rightarrow average 0.5, but F1 = 0)

1.4 fpr

- 1. The false positive rate (**FPR**) is the probability of a false alarm




1.5 roc_auc

- 1. The **area under the curve** of the **ROC**-curve, a measure between 0.5 (worthless) and 1.0 (perfect: getting no FPs), rates the ability of a model to achieve a blend of recall and precision

2.	3.	4.
-	-	-

The receiver operating characteristic (ROC) curve plots the TPR versus FPR as a function of the model's threshold for classifying a positive




Increasing the threshold → moving up a curve (\equiv model) to the top right corner, where all data is predicted as positive (threshold = 1.0) and vice versa

1.6 pr_auc

(see: [link 1](#) and [link 2](#))

1. The **area under the curve** of the **PR**-curve, same: similar objective as the **roc_auc**, but PR curves are better than ROC curves if the populations are imbalanced


The PR-curve plots the Precision versus the Recall



Instead of the top left corner for the ROC-curve, here, we want to be in the top right corner for our classifier to be perfect


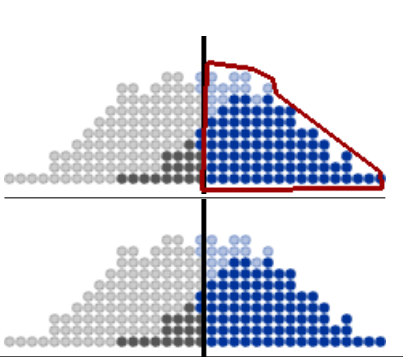
1.7 accuracy

1. Measuring the portion of correctly predicted data




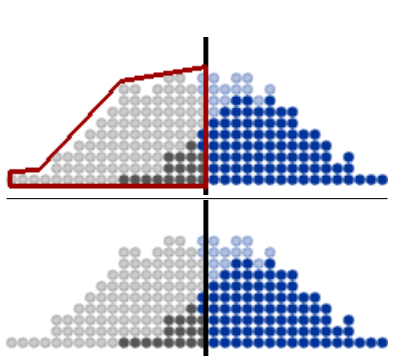
1.8 match_rate

1. The proportion of observations matched/not-matched

2. $= \frac{TP+FP}{Total}$	3. 	4. 
-----------------------------------	---	--

1.9 filter_rate

1. The proportion of observations filtered/not-filtered

2. $= 1 - \text{match_rate} = \frac{TN+FN}{Total}$	3. 	4. 
--	---	--

1.10 counts

1.

2.	3.	4.

1.11 rates

1.

2.	3.	4.

1.12 !<metric>

- Any <metric> with an exclamation mark is the same metric for the negative class
- e.g. $\text{recall} = \frac{TP}{TP+FN} \Rightarrow !\text{recall} = \frac{TN}{TN+FP}$
- Example usage: find all items that are not “E” class \rightarrow look at !recall for “E” class.


1.12.1 Existing !<metric>s

- !f1
- !precision
- !recall

1.13 Additional explanations

1.13.1 recall vs precision

When increasing one of these two, the other one naturally decreases. For an intuitive example, let's take a look at Google's Loan Threshold Simulation:



The dark grey / dark blue dots, representing clients that would actually pay back their loan, are more and more included (\rightarrow given loans) if we move the threshold further to the left.

But so are clients that would not. Thus moving the threshold to the left increases the **recall (tpr)** but decreases the **precision** and vice versa when moving to the right.

1.13.2 roc_auc vs pr_auc

see: <https://www.kaggle.com/general/7517>

- tl;dr: if the class imbalance problem exists, **pr_auc** is more appropriate than **roc_auc**

If TNs are not meaningful to the problem or there are a lot more negatives than positives, **pr_auc** is the way to go (it does not account for TNs).

- In other words:
 - If the model needs to perform equally on the positive and negative class \rightarrow **roc_auc**
 - If it's not interesting how the model performs on negative class \rightarrow **pr_auc** (example: detecting cancer; find all positives and make sure they're correct!)

Additional Information

- ORES Threshold Link
 - Confluence Link (Bachelor Thesis)
 - Amir's mail links
 - New Filters for Edit Review Documentation Link
 - ORES API Call Link
- “Basically it asks for threshold from the API when ”recall is at its maximum when precision is at least 0.995“
- Bowen Yu: Applying Value-Sensitive Algorithm Design to ORES

Questions

- **Q:** Should I ask Aaron how he would like us to work together? I'm not sure how he meant it.
A:
- **Q:** In what situations exactly do we want to optimize the threshold in the context of user centered threshold optimization?
A:
- **Q:** VPN recommendation?
A: