

ORES Custom Documentation V

Disclaimer: No guarantee for the correctness of information / explanations / sources is given.

Goals

1. Metric list:
 - Add examples
 - Adjust the description of counts \rightarrow labels, predictions
2. Research what form of revision data is needed (for existing visualizations, but also in general) \rightarrow also check RecentChanges and new Filters and what API calls look like in that context
3. Watch “ROC curves and Area Under the Curve explained” and think about what parameters could be used in which ways to filter the output of the current UI (currently: X inputs \rightarrow X outputs)

Also check out: Precision-Recall AUC vs ROC AUC discussion
4. Read A Review of User Interface Design for Interactive Machine Learning (carefully!)
5. Think about what could be the goal of this thesis

1 Crucial metrics: damaging-model

Metrics simple list:

!f1	✓
!precision	✓
!recall	✓
accuracy	✓
counts	✓
f1	✓
filter_rate	✓
fpr	✓
match_rate	✓
pr_auc	✓
precision	✓
rates	✓
recall	✓
roc_auc	✓

Again, changes have been made to the list and explanations, compared to the version in oresDoc4. The structure of explanations will be as follows:

For each metric (if possible) there will be:

1. The formula based on the **confusion matrix**
2. A definition
3. An intuitive explanation with an example
4. Its meaning based on the **loan threshold** representation by Google (Link)
5. Additional information (if necessary)

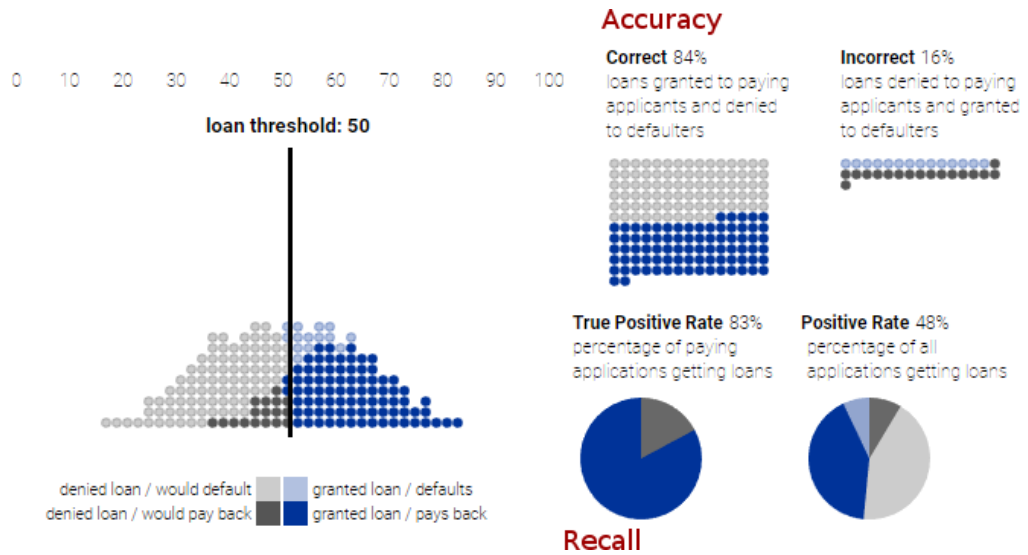
Explanations: References

- Confusion Matrix

		Actual	
		Positive	Negative
Predicted	Positive	True Positive	False Positive
	Negative	False Negative	True Negative

Abbreviations: **TP**, **FP**, **FN** and **TN**.

- Loan Threshold




The example scenario

To ease the understanding, let's stick to the following scenario and refer to it for each metric:

- 100 people represent our total population
- 35% of our population is infected with disease X

That leaves us with the following labels: 35 **positives** and 65 **negatives** (being tested for disease X;)




- A classifier is now supposed to classify every person of our population (based on their visible symptoms for example). This is the **prediction**:
 - If our algorithm says a person is infected, that person will be classified as a **positive**, and marked with radioactive symbol:



- If the prediction results in a **negative**, the person will be marked with a sun symbol:



- The classifier may predict that:
 - out of the 35 **infected people**, **30** are infected (those 30 are what we call **true positives**) and **5** are not (those 5 are **false negatives**)
 - out of the 65 **non infected people**, **10** are infected (**false positives**) and **55** are not (**true negatives**)



true positive: infected and correctly predicted (30)



false negative: infected and incorrectly predicted (5)





true negative: not infected and correctly predicted (55)



false positive: not infected and incorrectly predicted (10)

Let's get started.


1.1 recall

1. $\frac{TP}{TP+FN}$
2. Recall (\equiv true positive rate \equiv "sensitivity") is the ability of a model to find **all** relevant cases within the dataset.
3. Now, in our example scenario, the relevant cases are the infected people. We absolutely want to identify those: . The ability of the model to identify those depends on how many will be **correctly** predicted to be infected: .

In other words, we are looking for the ratio of correctly predicted to be infected people to all infected people.

That leads to $\frac{\sum \text{green robot}}{\sum \text{green robot}}$, with $\text{green robot} = \text{green robot} + \text{yellow robot}$, which is equivalent to the formula in 1., if you replace the symbols with their confusion matrix counterpart according to the legend in **The example scenario**.

In terms of numbers for our example that would be $\frac{30}{30+5} \approx 0.86$




1.2 precision

1. $\frac{TP}{TP+FP}$
2. Ability of the model to find **only** relevant cases within the dataset

- Again, we take a look at the relevant cases, the infected people: 🧑. This time around though, we are **not** interested in the ratio of correctly predicted to be infected people to **all** infected people. Instead we want to know how good the model is at only predicting those to be infected, that actually are. Therefore, we want the ratio of all 🧑 to all those predicted to be infected: 🧑 + 🧑

$$= \frac{\sum \text{🧑}}{\sum \text{🧑} + \sum \text{🧑}} = \frac{30}{30+10} = 0.75$$



1.3 f1


-
- f1-Score, the harmonic mean of recall and precision, a metric from **0** (worst) to **1** (best), used to evaluate the accuracy of a model by taking into account recall and precision: $= 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$
- For our example model, that would result in $= 2 * \frac{0.75 * \frac{30}{35}}{0.75 + \frac{30}{35}} = 0.8$
-
- Additional information: Compared to the simple average (of recall and precision), the harmonic mean punishes extreme values (e.g. precision 1.0 and recall 0.0 \rightarrow average = 0.5, but f1 = 0)

1.4 fpr

- $\frac{\text{FP}}{\text{FP} + \text{TN}}$

2. The false positive rate is the probability of a false alarm.
3. In our example, a false alarm would obviously be labeling someone as infected, who isn't: 🧑. Now we just have to ask ourselves what portion of those, that, if they **were** incorrectly predicted as infected (because they **are not** infected: 🧑), **are** incorrectly predicted as infected? Hence, we are looking for the ratio of 🧑 to all non infected people: 🧑 + 🧑.

$$= \frac{\sum \text{🧑}}{\sum \text{🧑} + \sum \text{🧑}} = \frac{10}{10+55} \approx 0.15$$




1.5 roc_auc

1. -
2. The **area under the curve** of the **ROC**-curve, a measure between 0.5 (worthless) and 1.0 (perfect: getting no FPs), rates the ability of a model to achieve a blend of recall and precision

2.	3.	4.
-	-	-

The receiver operating characteristic (ROC) curve plots the TPR versus FPR as a function of the model's threshold for classifying a positive




Increasing the threshold \rightarrow moving up a curve (\equiv model) to the top right corner, where all data is predicted as positive (threshold = 1.0) and vice versa

1.6 pr_auc

(see: link 1 and link 2)

1. The **area under the curve** of the **PR**-curve, same: similar objective as the **roc_auc**, but PR curves are better than ROC curves if the populations are imbalanced

The PR-curve plots the Precision versus the Recall




Instead of the top left corner for the ROC-curve, here, we want to be in the top right corner for our classifier to be perfect

1.7 accuracy

1. $\frac{TP+TN}{Total}$
2. Accuracy measures the portion of correctly predicted data
3. In our example scenario, this is equal to asking ourselves *out of all patients, what's the portion of correctly predicted cases?* The correctly predicted cases are infected patients, predicted to be infected (🧑🏻), and non infected patients, predicted not to be infected (🧑🏻). This wanted proportion results in:


$$= \frac{\sum \text{🧑🏻} + \sum \text{🧑🏻}}{\sum \text{🧑🏻} + \sum \text{🧑🏻}} = \frac{30+55}{35+65} = 0.85$$



1.8 match_rate

1. $\frac{TP+FP}{Total}$
2. The match rate is the proportion of observations matched/not-matched, meaning the ratio of observations predicted to be positive to the total of observations.
3. Concerning our example, this would be equal to wanting to know what portion of the population was predicted to be infected. Those groups are: 🧑🏻 and 🧑🏻.


$$= \frac{\sum \text{🧑🏻} + \sum \text{🧑🏻}}{\sum \text{🧑🏻} + \sum \text{🧑🏻}} = \frac{30+10}{35+65} = 0.4$$



1.9 filter_rate









1. $1 - \text{match_rate} = \frac{\text{TN} + \text{FN}}{\text{Total}}$
2. The filter rate is the proportion of observations filtered/not-filtered, meaning the ratio of observations predicted to be negative to the total of observations. This is the complement to the match rate.
3. In our example scenario, this would be equal to wanting to know what portion of the population was predicted not to be infected.

$$= \frac{\sum \text{👤} + \sum \text{👤}}{\sum \text{👤} + \sum \text{👤}} = \frac{55+5}{35+65} = 0.6 = 1 - \text{match_rate}$$



1.10 counts

1.
 - labels:
 - false: TN + FP
 - true: TP + FN
 - n: Total
 - predictions:
 - false:

- * false: TN
 - * true: FP
 - true:
 - * false: FN
 - * true: TP
2. • **labels:** The number of edits (*manually*) labeled as *false* and *true*: these values represent the **actual** positives and negatives.
- **n:** The sample size; total number of edits taken into account
 - **predictions:** edits ...
 - false: ... actually being false ...
 - * false: ... and predicted to be false
 - * true: ... but predicted to be true
 - true: ... actually being true ...
 - * false: ... but predicted to be false
 - * true: ... and predicted to be true
3. Concerning our example:
- **labels:**
 - false (**non infected people**):  +  = 55 + 10 = 65
 - true (**infected people**):  +  = 30 + 5 = 35
 - **n** (Total): 100
 - **predictions:**
 - false (**non infected people ...**)
 - * false (... and predicted not to be infected):  = 55
 - * true (... but predicted to be infected):  = 10
 - true (**infected people ...**)
 - * false (... predicted not to be infected):  = 5
 - * true (... predicted to be infected):  = 30
4. -
5. Additional information:

When calling the enwiki damaging model for example (Link), you will get the following output for counts:

```
"counts": {
  "labels": {
    "false": 18677,
    "true": 751
  },
  "n": 19428,
  "predictions": {
    "false": {
      "false": 17958,
      "true": 719
    },
    "true": {
      "false": 320,
      "true": 431
    }
  }
},
```

⇒ e.g. out of 18677 edits that were labeled as *false*, 719 were false positives

1.11 rates

1.
 - false: $\frac{\text{counts_labels_false}}{\text{counts_n}}$
 - true: $\frac{\text{counts_labels_true}}{\text{counts_n}}$
2. The rates simply give us the proportion of edits labeled as false or true to the total number of edits taken into account.
3. For the example scenario:
 - false (proportion of infected people to the total number of people tested): $\sum_{100}^{\text{👤}} = \frac{65}{100} = 0.65$
 - true (proportion of infected people to the total number of people tested): $\sum_{100}^{\text{👤}} = \frac{35}{100} = 0.35$
4. -
5. Additional information:

Calling the API the same way as for **counts** (Link), we get:

```

"rates": {
  "population": {
    "false": 0.966,
    "true": 0.034
  },
  "sample": {
    "false": 0.961,
    "true": 0.039
  }
},

```

The number of edits taken into account for **sample** equals the **n** from **counts** = 19428.

Now we see that, also looking at the output under **counts**, *rates_sample_false*:

$$0.961 = \frac{18677}{19428}.$$

Note that we are shown results for “population” and “sample”. There is a significant number of bot edits and edits that don’t need reviewing (admins, autopatrolled users). The **sample** of edits does not contain any of those.

1.12 !<metric>

- Any <metric> with an exclamation mark is the same metric for the negative class
- e.g. $\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \Rightarrow \text{!recall} = \frac{\text{TN}}{\text{TN} + \text{FP}}$
- Example usage: find all items that are not “E” class → look at !recall for “E” class.


1.12.1 Existing !<metric>s

- !f1
- !precision
- !recall

1.13 Additional explanations

1.13.1 recall vs precision

When increasing one of these two, the other one naturally decreases. For an intuitive example, let’s take a look at Google’s Loan Threshold Simulation:



The dark grey / dark blue dots, representing clients that would actually pay back their loan, are more and more included (\rightarrow given loans) if we move the threshold further to the left.

But so are clients that would not. Thus moving the threshold to the left increases the **recall (tpr)** but decreases the **precision** and vice versa when moving to the right.

1.13.2 roc_auc vs pr_auc

see: <https://www.kaggle.com/general/7517>

- tl;dr: if the class imbalance problem exists, **pr_auc** is more appropriate than **roc_auc**

If TNs are not meaningful to the problem or there are a lot more negatives than positives, **pr_auc** is the way to go (it does not account for TNs).

- In other words:
 - If the model needs to perform equally on the positive and negative class \rightarrow **roc_auc**

- If it's not interesting how the model performs on negative class
→ **pr_auc** (example: detecting cancer; find all positives and make sure they're correct!)

Questions

- **Q:** Should I ask Aaron how he would like us to work together? I'm not sure how he meant it.

A:

- **Q:** In what situations exactly do we want to optimize the threshold in the context of user centered threshold optimization?

A: